



CHARMe node Installation Document



CHARMe is funded by the EC under its FP7 Research Programme

Document Control

Contributors

Person	Role	Organisation	Contribution
A.Wilson	Developer	STFC	Initial Draft.

Document Approval

Person	Role	Organisation

References

ID	Author	Document Title	Date
[R-1]			

Revision History

Issue	Author	Date	Description
0.1	A.Wilson	22 nd Jan 2015	Initial Draft.
0.2	A.Wilson	17 th Feb 2015	Added Strabon instructions

Table of Contents

1 About these Instructions.....	4
2 Dependencies.....	4
3 Web Server.....	5
3.1 Certificate.....	5
3.2 Security.....	5
3.3 Additional Configuration for Fuseki.....	5
3.4 Start Up.....	5
4 Triple Store.....	6
4.1 Fuseki.....	6
4.1.1 Installation.....	6
4.1.2 Configuration.....	6
5 CHARMe.....	9
5.1 Installation.....	9
5.2 Configuration.....	9
5.2.1 local_settings.py.....	9
5.2.2 wsgi.py.....	10
5.2.3 djcharme_wsgi.conf.....	10
5.2.4 Other Configuration.....	10
5.2.5 Set up the Client via the GUI.....	11
5.2.6 oauth_test2.html.....	11
6 Strabon.....	12

Node Services

The CHARMe node is a Python Django application that sits behind an Apache web server.

1 About these Instructions

The instructions in this document have been tested on Redhat Enterprise Linux and relate to the installation of the central CHARMe node at STFC. When installing on a different node the text in **green** will need to be changed. The text in **red** WILL need to be changed for all installations.

2 Dependencies

The server has dependencies on java, python-virtualenv, gcc, mod_ssl and mod_wsgi. These can be installed with:

```
yum install java-1.7.0-openjdk.x86_64 python-virtualenv.noarch gcc mod_ssl mod_wsgi
```

It is possible to use one of a number of databases for storing the admin data. For a production node it is recommended that postgres is used.

3 Web Server

An Apache web server is used as a front end for Fuseki and CHARMe.

3.1 Certificate

Apache should be configured to use a certificate, such as a Comodo certificate.

3.2 Security

There are a number of things to do to tighten security. Turn off TRACE, in **httpd.conf** add the following:

```
TraceEnable off
```

Disable SSL v2 and low level ciphers, in **ssl.conf**:

```
SSLProtocol all -SSLv2 -SSLv3
SSLCipherSuite ALL:!ADH:!EXPORT:!LOW:!SSLv2:!RC4:+RSA:+HIGH:+MEDIUM
```

3.3 Additional Configuration For Fuseki

In order to give access to the Fuseki query page and end point, add the following to **httpd.conf**:

```
ProxyPass /sparql.html http://127.0.0.1:3333/sparql.html
ProxyPassReverse /sparql.html http://127.0.0.1:3333/sparql.html
ProxyPass /sparql http://127.0.0.1:3333/privateds/sparql
ProxyPassReverse /sparql http://127.0.0.1:3333/privateds/sparql
ProxyPass /fuseki.css http://127.0.0.1:3333/fuseki.css
ProxyPassReverse /fuseki.css http://127.0.0.1:3333/fuseki.css
```

3.4 Start Up

Start the service:

```
/etc/init.d/httpd start
chkconfig httpd on
```

4 Triple Store

The CHARMe node has been tested with Fuseki and Strabon.

4.1 Fuseki

4.1.1 Installation

Create the required directories:

```
mkdir -p /opt/charme/luceneDB /var/log/fuseki
chown apache /opt/charme
```

Get the latest version of Fuseki and unpack it:

```
cd /opt/
wget http://mirror.vorboss.net/apache/jena/binaries/jena-fuseki-1.1.1-
distribution.zip
unzip jena-fuseki-1.1.1-distribution.zip
mv jena-fuseki-1.1.1 jena-fuseki
```

4.1.2 Configuration

Set up the start up script:

```
cp /opt/jena-fuseki/fuseki /etc/init.d/
```

In **/etc/init.d/fuseki** add the FUSEKI_* values and edit JAVA_OPTIONS:

```
export FUSEKI_HOME="/opt/jena-fuseki"
export FUSEKI_ARGS="--update --port=3333 --config=/opt/charme/config-charme.ttl"
export FUSEKI_DATA_DIR="/opt/charme/DB"
export FUSEKI_LOGS="/var/log/fuseki"
JAVA_OPTIONS="-Dlog4j.configuration=file:/opt/jena-fuseki/log4j.properties
-Xmx1200M"
```

Create the file **/opt/charme/config-charme.ttl** with the contents:

```
# Licensed under the terms of http://www.apache.org/licenses/LICENSE-2.0

## Example of a TDB dataset published using Fuseki: persistent storage.

@prefix :      <#> .
@prefix fuseki: <http://jena.apache.org/fuseki#> .
@prefix rdf:    <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .
@prefix rdfs:   <http://www.w3.org/2000/01/rdf-schema#> .
@prefix tdb:    <http://jena.hpl.hp.com/2008/tdb#> .
@prefix ja:     <http://jena.hpl.hp.com/2005/11/Assembler#> .
@prefix text:   <http://jena.apache.org/text#> .
@prefix dcterms: <http://purl.org/dc/terms/> .

[] rdf:type fuseki:Server ;
  # Timeout - server-wide default: milliseconds.
  # Format 1: "1000" -- 1 second timeout
  # Format 2: "10000,60000" -- 10s timeout to first result, then 60s timeout to for
rest of query.
  # See java doc for ARQ.queryTimeout
  # ja:context [ ja:cxtName "arq:queryTimeout" ; ja:cxtValue "10000" ] ;
  # ja:loadClass "your.code.Class" ;

  fuseki:services (
    <#service_tdb_all>
    <#service_text_tdb>
  ) .

# TDB
[] ja:loadClass "com.hp.hpl.jena.tdb.TDB" .
tdb:DatasetTDB rdfs:subClassOf ja:RDFDataset .
tdb:GraphTDB   rdfs:subClassOf ja:Model .

## Initialize text query
[] ja:loadClass "org.apache.jena.query.text.TextQuery" .
# A TextDataset is a regular dataset with a text index.
text:TextDataset rdfs:subClassOf ja:RDFDataset .
# Lucene index
text:TextIndexLucene rdfs:subClassOf text:TextIndex .

# Text index description
<#indexLucene> a text:TextIndexLucene ;
  text:directory <file:/opt/charme/luceneDB> ;
  text:entityMap <#entMap> ;
  .

# Mapping in the index
# URI stored in field "uri"
# rdfs:label is mapped to field "text"
<#entMap> a text:EntityMap ;
  text:entityField "uri" ;
  text:defaultField "title" ;
  text:field "creator" ;
  text:map (
    [ text:field "title" ; text:predicate dcterms:title ]
    [ text:field "creator" ; text:predicate dcterms:creator ]
  ) .
```

```

## -----
### This URI must be fixed - it's used to assemble the text dataset.
#
:text_dataset rdf:type      text:TextDataset ;
    text:dataset    <#tdb_dataset_readwrite> ;
    text:index      <#indexLucene> ;
    .

<#service_text_tdb> rdf:type fuseki:Service ;
    rdfs:label      "TDB Service (full_text)" ;
    fuseki:name      "privateds" ;
    fuseki:serviceQuery      "query" ;
    fuseki:serviceQuery      "sparql" ;
    fuseki:serviceUpdate      "update" ;
    fuseki:serviceUpload      "upload" ;
    fuseki:serviceReadWriteGraphStore "data" ;
    # A separate read-only graph store endpoint:
    fuseki:serviceReadGraphStore      "get" ;
    fuseki:dataset      :text_dataset ;
    .

## -----
## Updatable TDB dataset with all services enabled.

<#service_tdb_all> rdf:type fuseki:Service ;
    rdfs:label      "TDB Service (RW)" ;
    fuseki:name      "privateds" ;
    fuseki:serviceQuery      "query" ;
    fuseki:serviceQuery      "sparql" ;
    fuseki:serviceUpdate      "update" ;
    fuseki:serviceUpload      "upload" ;
    fuseki:serviceReadWriteGraphStore "data" ;
    # A separate read-only graph store endpoint:
    fuseki:serviceReadGraphStore      "get" ;
    fuseki:dataset      <#tdb_dataset_readwrite> ;
    .

<#tdb_dataset_readwrite> rdf:type      tdb:DatasetTDB ;
    tdb:location "/opt/charme/DB" ;
    tdb:unionDefaultGraph true ;
##      # Query timeout on this dataset (milliseconds)
##      ja:context [ ja:cxtName "arq:queryTimeout" ; ja:cxtValue "1000" ] ;
##      # Default graph for query is the (read-only) union of all named graphs.

```

Start the services:

```

/etc/init.d/fuseki start
chkconfig fuseki on
chown apache:apache /opt/charme/charme.db
/etc/init.d/httpd restart

```


5 CHARMe

5.1 Installation

Set up the python virtual environment:

```
cd /opt/  
virtualenv djcharme  
source djcharme/bin/activate  
cd djcharme/  
export http_proxy=http://wwwcache.rl.ac.uk:8080  
export https_proxy=http://wwwcache.rl.ac.uk:8080  
mkdir /var/www/html/djcharme  
export DJANGO_PROJECT_STATIC_FILES=/var/www/html/djcharme/
```

Set the version of the code you wish to install:

```
export VERSION=0.7.2
```

Get the code:

```
wget https://github.com/cedadev/djcharme/raw/develop/djcharme/dist/djcharme-${VERSION}.tar.gz
```

Install the postgres python libraries and the CHARMe code:

```
pip install psycopg2  
pip install djcharme-${VERSION}.tar.gz --extra-index-url http://dist.ceda.ac.uk/pip/
```

5.2 Configuration

5.2.1 local_settings.py

This shows the configuration required for using postgres as the admin database. Edit **/opt/djcharme/lib/python2.6/site-packages/djcharme/local_settings.py**.

In the DATABASES section set values for ENGINE, NAME, USER and PASSWORD, e.g.

```
'ENGINE': 'django.db.backends.postgresql_psycopg2',  
'NAME': 'charmedb',  
'USER': 'postgresUser',  
'PASSWORD': 'myPassword',
```

Also set values for `STATIC_ROOT` and `NODE_URI`, e.g.

```
STATIC_ROOT = '/var/www/html/djcharme/'  
NODE_URI = 'https://charme.cems.rl.ac.uk'
```

If needed, set the values for `HTTP_PROXY` and `HTTP_PROXY_PORT`.

Update the email settings as required. `DEFAULT_FROM_EMAIL` will need to be set, e.g.

```
DEFAULT_FROM_EMAIL = 'no-reply@charme.cems.rl.ac.uk'
```

5.2.2 wsgi.py

Update **wsgi.py**:

```
sed -i s+VEPATH_PAR+/opt/djcharme/lib/python2.6/site-packages+  
/opt/djcharme/lib/python2.6/site-packages/djcharme/resources/wsgi.py  
sed -i s+DJANGO_PATH+/opt/djcharme/lib/python2.6/site-packages/djcharme+  
/opt/djcharme/lib/python2.6/site-packages/djcharme/resources/wsgi.py  
sed -i s+PROJECT_LIB_PATH+/opt/djcharme/lib/python2.6/site-packages/djcharme+  
/opt/djcharme/lib/python2.6/site-packages/djcharme/resources/wsgi.py
```

5.2.3 djcharme_wsgi.conf

Set up the **djcharme_wsgi.conf**:

```
cp /opt/djcharme/lib/python2.6/site-  
packages/djcharme/resources/djcharme_wsgi.conf /etc/httpd/conf.d/  
sed -i s+LOG_DIR_PATH_PAR/PROJECT_NAME_PAR+/etc/httpd/logs/djcharme+  
/etc/httpd/conf.d/djcharme_wsgi.conf  
sed -i s+PATH_TO_DJANGO_PROJECT_STATIC_FILES+/var/www/html/djcharme/+  
/etc/httpd/conf.d/djcharme_wsgi.conf  
sed -i s+PATH_TO_PROJECT_WSGI+/opt/djcharme/lib/python2.6/site-  
packages/djcharme/resources+ /etc/httpd/conf.d/djcharme_wsgi.conf  
sed -i s+PROJECT_NAME_PAR++ /etc/httpd/conf.d/djcharme_wsgi.conf
```

5.2.4 Other Configuration

Initialise the database:

```
python /opt/djcharme/lib/python2.6/site-packages/djcharme/manage.py collectstatic  
--clear --noinput  
python /opt/djcharme/lib/python2.6/site-packages/djcharme/manage.py syncdb --noinput
```

Set up the admin user:

```
python /opt/djcharme/lib/python2.6/site-packages/djcharme/manage.py createsuperuser
```

Deactivate the python environment:

```
deactivate
```

Restart the service:

```
/etc/init.d/httpd restart
```

5.2.5 Set up the Client via the GUI

In a browser go to the admin interface:

<https://charme.cems.rl.ac.uk/admin/>

Then go to clients, Add client and provide the relevant values, e.g.

Url: <https://charme.cems.rl.ac.uk/>

Redirect uri: <https://charme.cems.rl.ac.uk/>

Client type: Public (Native and JS applications)

Organization: STFC

Make a note of the Client id.

5.2.6 oauth_test2.html

Edit /opt/djcharme/lib/python2.6/site-packages/djcharme/templates/oauth_test2.html.

Set the values of oa_domain and client_id, using the client id from above, e.g.

```
oa_domain='https://charme-test.cems.rl.ac.uk'  
client_id='1234567890'
```

Restart the service:

```
/etc/init.d/httpd restart
```

6 Strabon

Only very basic testing has been undertaken with the Strabon triple store. To use Strabon in place of Fuseki a number of changes are needed in the file `local_settings.py`, i.e.:

```
SPARQL_PORT = "8080"
SPARQL_DATASET = "strabonendpoint"
SPARQL_UPDATE = _format_sparql_url("Update")
SPARQL_QUERY = _format_sparql_url("Query")
```

In addition to these changes three extra properties are required:

```
SPARQL_USERNAME = "endpoint"
SPARQL_PASSWORD = "myPassword"
STRABON = True
```

This should be all of the extra configuration required to run the CHARMe node using Strabon.